


# Artificial Intelligence Needs **Meta Intelligence**

The Case for Metacognitive AI — self-monitoring plus resource-rational control as a design principle

Sergei Chuprov<sup>1</sup>, Richard D. Lange<sup>2</sup>, Leon Reznik<sup>2</sup>, Paulo Shakarian<sup>3</sup>, Raman Zatsarenko<sup>2</sup>, Dmitrii Korobeinikov<sup>2</sup>

<sup>1</sup>University of Texas Rio Grande Valley · <sup>2</sup>Rochester Institute of Technology · <sup>3</sup>Syracuse University

✉ sergei.chuprov@utrgv.edu · rdlvcs@rit.edu · leon.reznik@rit.edu · pashakar@syr.edu · rz4983@rit.edu · dk9148@rit.edu



## POSITION OF THIS PAPER

**Metacognition**, encompassing self-monitoring and resource-allocation inspired by cognitive science, should be a **general design principle for AI systems**. Further, we call for a **theoretical framework** to unify current efforts, and **actionable strategies** to translate metacognition research into practical AI.

## Why this matters now

Recent gains in inference and training come at a **soaring computational cost**. Reasoning LLMs and large-scale training strain time, memory, and energy, yet give few guarantees on correctness or robustness — and the way we decide how much computation a problem deserves stays crude, usually a fixed budget set by hand.

The same pressure appears everywhere compute is finite and stakes vary across inputs: from perception on the edge to multi-step reasoning in the cloud. Treating every input the same wastes effort on the easy ones and starves the hard ones.

**Humans tend to do better in this:** we *sense* when a task is hard, *deliberate* longer when it matters, and *stop* when it does not. A century of cognitive science calls this **metacognition**. A metacognitive system spends resources **up front**, per instance (Fig. 1).

The goal is not less computation, but **computation spent where it counts**.

## Objective, InteFL, case study

A resource-rational objective trades task loss against compute:


$$\min_{w,w'} \frac{1}{N} \sum_{i=1}^N \mathbb{E}_{c_i \sim \pi_{w'}} [\text{loss}(x_i, f_w(x_i; c_i)) + \alpha \text{cost}(c_i)] + \beta \|\pi_{w'}\|$$

Federated Learning (FL) is naturally well-positioned for metacognition: it already separates an object level (local client training) from a meta level (server aggregation). A monitoring function  $\mathcal{M}$  scores each update by trust and distance from consensus; control accepts, rejects, or re-weights it (Fig. 2), self-adapting under drift or attack.

In **InteFL** (Fig. 3) [1], the aggregation strategy is the control mode  $c$ : the object-level step is one aggregation over client updates,  $\pi$  selects among strategies (FedAvg [6], Multi-Krum, PID-based exclusion), and its parameters  $w'$  are the exclusion thresholds tuned against observed performance. In **learning**, investing in a controller is especially worthwhile because its overhead is **amortized over the model's lifetime**.

On OctMNIST (Fig. 4), a tuned PID rule (2.5 $\sigma$ ) converges lowest and best excludes malicious clients; Other baselines miss attackers in most cases.

**Figure 4.** OctMNIST, 20 clients (4 adversarial): the tuned PID rule converges lowest; Trimmed Mean and Bulyan miss attackers and stall.



## What metacognition is


Following Nelson and Narens [2], cognition is **bi-level**: object-level processes (perception, learning, reasoning) are watched by a **meta-level**.

It adds **monitoring** — reading one's own state via cheap cues — and **control**, reallocating resources or switching strategy: a System-2 intervention [3].

## Research tasks

We call for concerted progress on three research tasks:

- 1. Characterize the performance–compute trade-off** — relate accuracy, robustness, and generalization to time, memory, and energy (PAC learnability, anytime algorithms, efficient models).
  - 2. Develop robust, low-overhead cues** — error-detection rules, early-exit signals, XAI attribution — reliable yet computable in near-constant time.
  - 3. Build well-calibrated controllers** with uncertainty quantification (temperature scaling, conformal prediction, abduction) [4].
- Solving these turns scattered tricks into **cumulative science** and provably good metacognitive control.




## Challenges and alternatives

**Alternatives.** “Scale is all you need” [5] and post-hoc “generate-and-verify” pipelines raise accuracy but not efficiency or real-time robustness; and unlike early-exit or mixture-of-experts, a metacognitive controller is an explicit, learned, runtime-active layer.

**Challenges** — three open problems:

- 1. Cues and controllers at low overhead** (compute and system complexity).
- 2. Translating cognitive-science concepts into precise computational infrastructure.**
- 3. Failures of metacognition:** detection may not be actionable [7], and controllers can misjudge competence or act globally-counterproductively.



## Conclusion

We argue that communities from cognitive science to AI and ML should jointly develop metacognition and make it a **leading paradigm of AI design**.

Long posited as a key enabler of general intelligence, metacognition is — on our pragmatic view — a concrete paradigm for managing the growing computational cost of AI, and a practical path to efficient, robust, and secure systems, as **InteFL** [1] demonstrates for FL.

## References

- [1] D. Korobeinikov, R. Zatsarenko, S. Chuprov, A. Barea, and L. Reznik. InteFL Framework: Optimizing Federated Learning with Metacognition for Application Design and Deployment. IEEE Intelligent Systems, pp. 1–10, 2026. DOI: 10.1109/MIS.2026.3658072.
- [2] T. O. Nelson and L. Narens. Metamemory: A Theoretical Framework and New Findings. The Psychology of Learning and Motivation, Vol. 26, pp. 125–173, Academic Press, 1990. DOI: 10.1016/S0079-7421(08)60053-5.
- [3] R. Ackerman and V. A. Thompson. Meta-Reasoning: Monitoring and Control of Thinking and Reasoning. Trends in Cognitive Sciences, Vol. 21, No. 8, pp. 607–617, 2017. DOI: 10.1016/j.tics.2017.05.004.
- [4] F. Lieder and T. L. Griffiths. Resource-Rational Analysis: Understanding Human Cognition as the Optimal Use of Limited Computational Resources. Behavioral and Brain Sciences, Vol. 43, Article e1, Cambridge University Press, 2020. DOI: 10.1017/S0140525X1900061X.
- [5] R. S. Sutton. The Bitter Lesson. Incomplete Ideas (Self-published essay), Mar. 13, 2019. Available: <http://www.incompleteideas.net/IncIdeas/BitterLesson.html>.
- [6] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. Agüera y Arcas. Communication-Efficient Learning of Deep Networks from Decentralized Data. Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS), PMLR Vol. 54, pp. 1273–1282, 2017. arXiv: 1602.05829.
- [7] P. Shakarian, G. I. Simari, and N. D. Bastian. Probabilistic Foundations for Metacognition via Hybrid-AI. arXiv preprint arXiv:2502.05398, 2025.

